

# CS 331: Artificial Intelligence

## Informed Search

1

## Informed Search

- How can we make search smarter?
- Use problem-specific knowledge beyond the definition of the problem itself
- Specifically, incorporate knowledge of how good a non-goal state is

2

## Best-First Search

- Node selected for expansion based on an evaluation function  $f(n)$ . I.e., expand the node that *appears* to be the best
- Node with lowest evaluation is selected for expansion
- Uses a priority queue
- We'll talk about Greedy Best-First Search and A\* Search

3

## Heuristic Function

- $h(n)$  = estimated cost of the cheapest path from node  $n$  to a goal node
- Non-negative
- $h(\text{goal node}) = 0$
- Contains additional knowledge of the problem

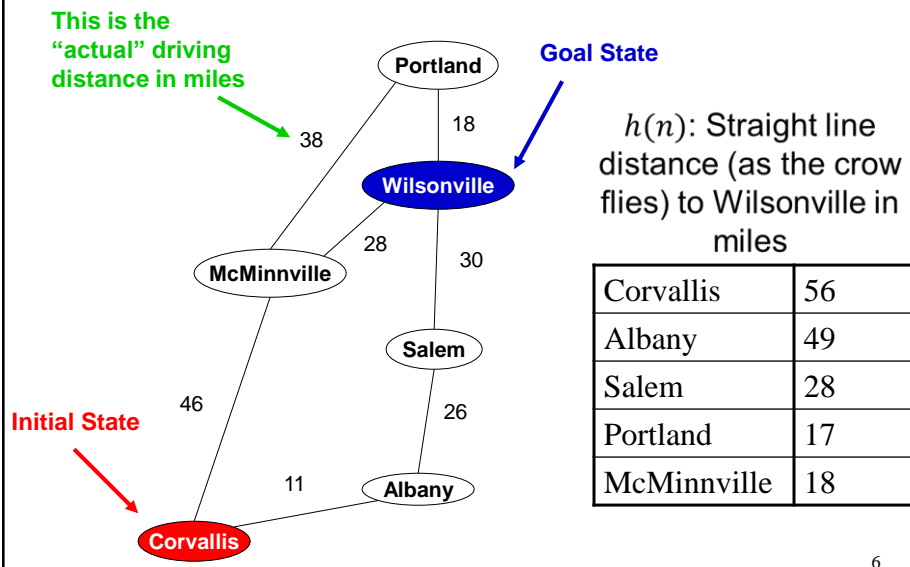
4

## Greedy Best-First Search


- Expands the node that is closest to the goal
- $f(n) = h(n)$

5

## Greedy Best-First Search Example



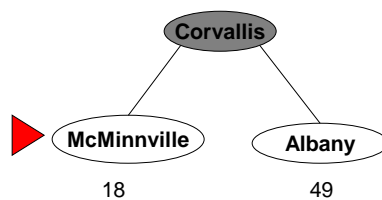
## Greedy Best-First Search Example



Corvallis	56
Albany	49
Salem	28
Portland	17
McMinnville	18

7

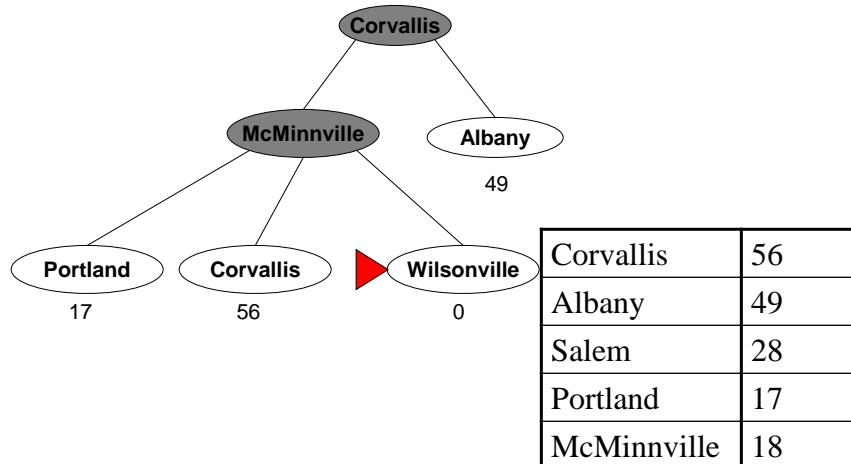
## Greedy Best-First Search Example



Corvallis	56
Albany	49
Salem	28
Portland	17
McMinnville	18

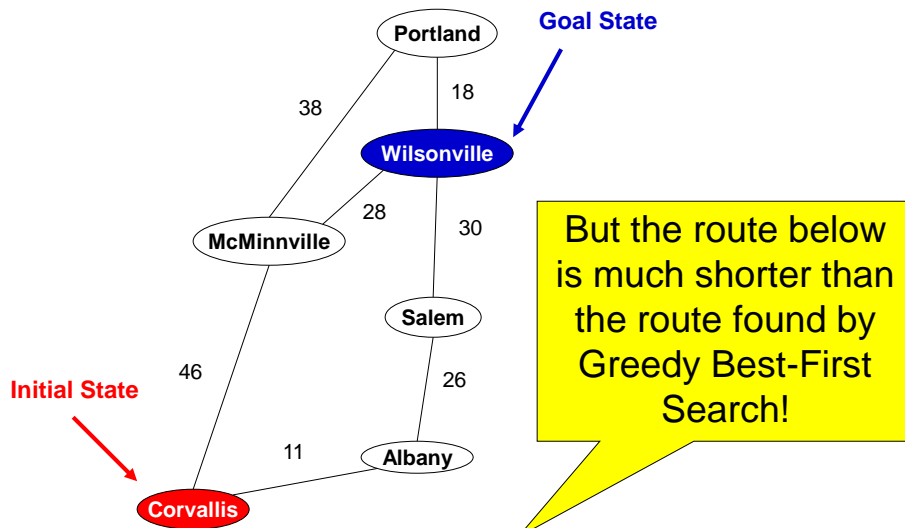
8

## Greedy Best-First Search Example



Corvallis → McMinnville → Wilsonville = 74 miles

## Greedy Best-First Search Example



Corvallis → Albany → Salem → Wilsonville = 67 miles

## Evaluating Greedy Best-First Search

Complete?	No (could start down an infinite path)
Optimal?	
Time Complexity	
Space Complexity	

11

## Evaluating Greedy Best-First Search

Complete?	No (could start down an infinite path)
Optimal?	No
Time Complexity	
Space Complexity	

12

## Evaluating Greedy Best-First Search

Complete?	No (could start down an infinite path)
Optimal?	No
Time Complexity	$O(b^m)$
Space Complexity	

13

## Evaluating Greedy Best-First Search

Complete?	No (could start down an infinite path)
Optimal?	No
Time Complexity	$O(b^m)$
Space Complexity	$O(b^m)$

Greedy Best-First search results in lots of dead ends which leads to unnecessary nodes being expanded

14

## A\* Search

- A much better alternative to greedy best-first search
- Evaluation function for A\* is:
$$f(n) = g(n) + h(n)$$
where  $g(n)$  = path cost from the start node to  $n$
- If  $h(n)$  satisfies certain conditions, A\* search is optimal and complete!

15

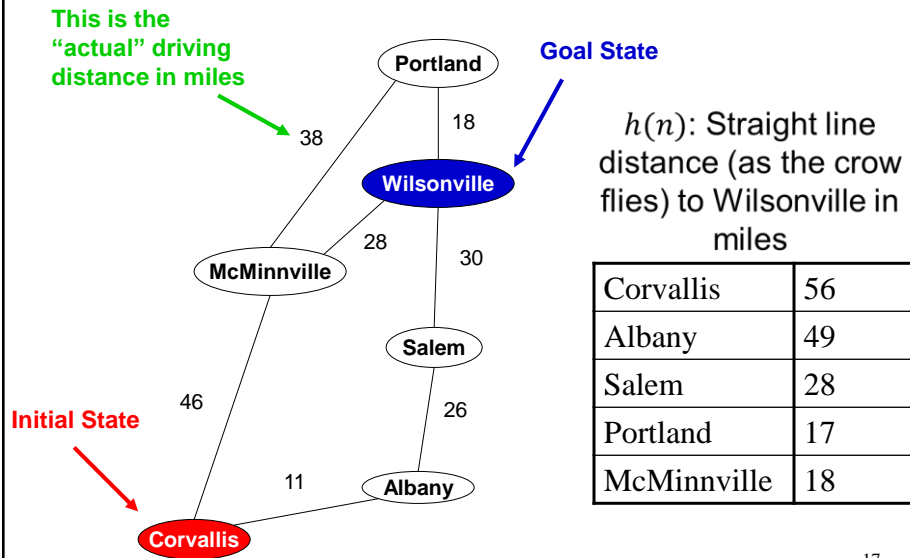
## Admissible Heuristics

- A\* is optimal if  $h(n)$  is an admissible heuristic
- An admissible heuristic is one that never overestimates the cost to reach the goal
- Admissible heuristic = optimistic
- Straight line distance was an admissible heuristic

16




## Greedy Best-First Search Example



17

## A\* Search Example




$h(n)$ : Straight line distance (as the crow flies) to Wilsonville in miles

Corvallis	56
Albany	49
Salem	28
Portland	17
McMinnville	18

18

## A\* Search Example




$h(n)$ : Straight line distance (as the crow flies) to Wilsonville in miles

Corvallis	56
Albany	49
Salem	28
Portland	17
McMinnville	18

19

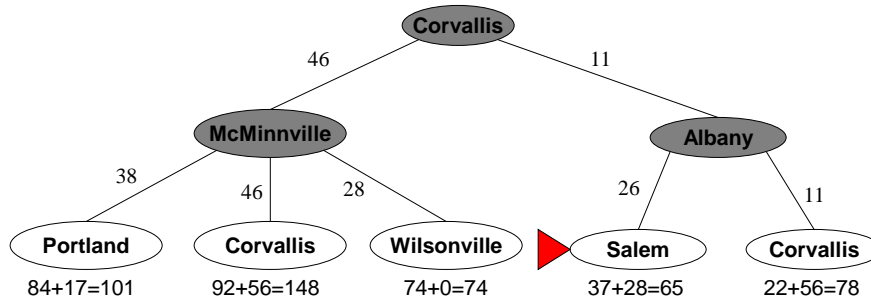
## A\* Search Example



Corvallis	56
Albany	49
Salem	28
Portland	17
McMinnville	18

20

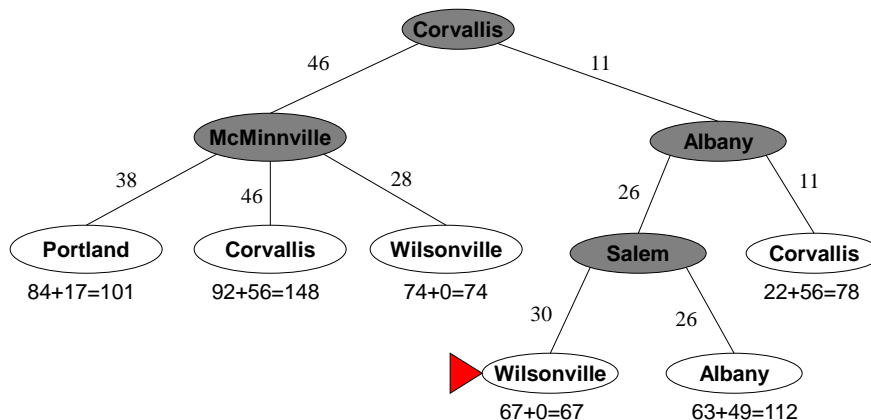
## A\* Search Example



Note: Don't stop when you put a goal state on the priority queue (otherwise you get a suboptimal solution)

21

## A\* Search Example



Proper termination: Stop when you pop a goal state from the priority queue


## Proof that A\* using TREE-SEARCH is optimal if $h(n)$ is admissible

- Suppose A\* returns a suboptimal goal node  $G_2$ .
- $G_2$  must be the least cost node in the frontier. Let the cost of the optimal solution be  $C^*$ .
- Because  $G_2$  is suboptimal:  

$$f(G_2) = g(G_2) + h(G_2) = g(G_2) > C^*$$
- Now consider a frontier node  $n$  on an optimal solution path to the goal  $G$ .
- If  $h(n)$  is admissible, then:  

$$f(n) = g(n) + h(n) \leq C^*$$
- We have shown that  $f(n) \leq C^* < f(G_2)$ , so  $G_2$  will not get expanded before  $n$ . Hence A\* must return an optimal solution.

$h(G_2) = 0$  because it is a goal node



23

## What about search graphs (more than one path to a node)?

- What if we expand a state we've already seen?
- Suppose we use the GRAPH-SEARCH solution and not expand repeated nodes
- Could discard the optimal path if it's not the first one generated
- One simple solution: ensure optimal path to any repeated state is always the first one followed (like in Uniform-cost search)
- Requires an extra requirement on  $h(n)$  called consistency (or monotonicity)


24

# Consistency

- A heuristic is **consistent** if, for every node  $n$  and every successor  $n'$  of  $n$  generated by any action  $a$ :

$$h(n) \leq \underbrace{c(n, a, n')}_{\text{Step cost of going from } n \text{ to } n' \text{ by doing action } a} + h(n')$$

- A form of the triangle inequality – each side of the triangle cannot be longer than the sum of the two sides



# Consistency

- Every consistent heuristic is also admissible**
- A\* using GRAPH-SEARCH is optimal if  $h(n)$  is consistent
- Most admissible heuristics are also consistent

## Consistency

- Claim: If  $h(n)$  is consistent, then the values of  $f(n)$  along any path are non-decreasing

- Proof:

Suppose  $n'$  is a successor of  $n$ . Want to show  $f(n') \geq f(n)$

$$g(n') = g(n) + c(n, a, n') \text{ for some } a$$

$$f(n') = g(n') + h(n')$$

$$= g(n) + c(n, a, n') + h(n')$$

$$\geq g(n) + h(n) \quad \leftarrow \text{From defn of consistency:}$$

$$= f(n) \quad c(n, a, n') + h(n') \geq h(n)$$

- Thus, the sequence of nodes expanded by A\* is in non-decreasing order of  $f(n)$
- First goal selected for expansion must be an optimal solution since all later nodes will be at least as expensive

27

## A\* is Optimally Efficient

- Among optimal algorithms that expand search paths from the root, A\* is **optimally efficient** for any given consistent heuristic function
- Optimally efficient: no other optimal algorithm is guaranteed to expand fewer nodes than A\*
  - Fine print: except A\* might possibly expand more nodes with  $f(n) = C^*$  where  $C^*$  is the cost of the optimal path – tie-breaking issues
- Any algorithm that does not expand all nodes with  $f(n) < C^*$  runs the risk of missing the optimal solution

28

## Evaluating A\* Search

With a consistent heuristic, A\* is complete, optimal and optimally efficient. Could this be the answer to our searching problems?

29

## Evaluating A\* Search

With a consistent heuristic, A\* is complete, optimal and optimally efficient. Could this be the answer to our searching problems?

### The Dark Side of A\*...



Time complexity is exponential (although it can be reduced significantly with a good heuristic)

The really bad news: space complexity is exponential (usually need to store all generated states). Typically runs out of space on large-scale problems.

30

## Summary of A\* (tree) Search

Complete?	Yes if $h(n)$ is admissible, $b$ is finite, and all step costs exceed some finite $\epsilon$
Optimal?	
Time Complexity	
Space Complexity	

## Summary of A\* (tree) Search

Complete?	Yes if $h(n)$ is admissible, $b$ is finite, and all step costs exceed some finite $\epsilon$
Optimal?	Yes if $h(n)$ is admissible
Time Complexity	
Space Complexity	



## Summary of A\* (tree) Search

Complete?	Yes if $h(n)$ is admissible, $b$ is finite, and all step costs exceed some finite $\epsilon$
Optimal?	Yes if $h(n)$ is admissible
Time Complexity	$O(b^d)$ (In the worst case but a good heuristic can reduce this significantly)
Space Complexity	

## Summary of A\* (tree) Search

Complete?	Yes if $h(n)$ is admissible, $b$ is finite, and all step costs exceed some finite $\epsilon$
Optimal?	Yes if $h(n)$ is admissible
Time Complexity	$O(b^d)$ (In the worst case but a good heuristic can reduce this significantly)
Space Complexity	$O(b^d)$ – Needs $O(\text{number of states})$ , will run out of memory for large search spaces

## Summary of A\* (graph) Search

Complete?	Yes if $h(n)$ is consistent, $b$ is finite, and all step costs exceed some finite $\epsilon$ <sup>1</sup>
Optimal?	Yes if $h(n)$ is consistent, $b$ is finite, and all step costs exceed some finite $\epsilon$ <sup>1</sup>
Time Complexity	$O(b^d)$ (In the worst case but a good heuristic can reduce this significantly)
Space Complexity	$O(b^d)$ – Needs $O(\text{number of states})$ , will run out of memory for large search spaces

<sup>1</sup> Since  $f(n)$  is non-decreasing, we must eventually hit an  $f(n) = \text{cost of the path to a goal state}$

## Iterative Deepening A\*

- Use iterative deepening trick to reduce memory requirements for A\*
  - In each iteration do a “cost-limited” depth first search.
    - Cutoff is based on the f-cost ( $g+h$ ) rather than the depth
  - After each iteration, the new cutoff is the smallest f-cost that exceeded the cutoff in the previous iteration
- Complete, Optimal but more costly than A\* and can take a while to run with real-valued costs

36

## Examples of heuristic functions

### The 8-puzzle

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8

End State

Heuristic #1:  $h_1$  = number of misplaced tiles e.g., start state has 8 misplaced tiles. This is an admissible heuristic.

37

## Examples of heuristic functions

### The 8-puzzle

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8

End State

Heuristic #2:  $h_2$  = total Manhattan distance (sum of horizontal and vertical moves, no diagonal moves). Start state is  $3+1+2+2+3+2+2+3=18$  moves away from the end state. This is also an admissible heuristic.

38

## Which heuristic is better?

- $h_2$  dominates  $h_1$ . That is, for any node  $n$ ,  $h_2(n) \geq h_1(n)$ .
- $h_2$  never expands more nodes than A\* using  $h_1$  (except possibly for some nodes with  $f(n) = C^*$ )

Proof:

Let  $h_1$  and  $h_2$  be admissible heuristics.

Every node with  $f(n) < C^*$  will surely be expanded, since A\* is optimal with an admissible heuristic. Since  $f(n) = g(n) + h(n)$ , every node with  $h(n) < C^* - g(n)$  will surely be expanded for either heuristic.

Since  $h_2$  is at least as big as  $h_1$  for all nodes, every node expanded with A\* using  $h_2$  will also be expanded with A\* using  $h_1$ . But  $h_1$  might expand other nodes as well. In other words, we have  $h_1(n) \leq h_2(n) < C^* - g(n)$

39

## Which heuristic is better?

Depth	# nodes expanded		
	IDS	A*( $h_1$ )	A*( $h_2$ )
2	10	6	<b>6</b>
4	112	13	<b>12</b>
6	680	20	<b>18</b>
8	6384	39	<b>25</b>
10	47127	93	<b>39</b>
12	3644035	227	<b>73</b>
14		539	<b>113</b>
16		1301	<b>211</b>
18		3056	<b>363</b>
20		7276	<b>676</b>
22		18094	<b>1219</b>
24		39135	<b>1641</b>

From Russell and Norvig Figure 4.8 (Results averaged over 100 instances of the 8-puzzle for depths 2-24).

## Inventing Admissible Heuristics

- Relaxed problem: a problem with fewer restrictions on the actions
- The cost of an optimal solution to a relaxed problem is an admissible heuristic for the original problem
- If we relax the rules so that a square can move anywhere, we get heuristic  $h_1$
- If we relax the rules to allow a square to move to any adjacent square, we get heuristic  $h_2$

41

## What you should know

- Be able to run A\* by hand on a simple example
- Why it is important for a heuristic to be admissible and consistent
- Pros and cons of A\*
- How do you come up with heuristics
- What it means for a heuristic function to dominate another heuristic function

42